

Importing the Dependencies

```
from google.colab import drive
drive.mount('/content/drive')

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

Data collection and Processing

```
#loading the csv data to Pandas
heart_data = pd.read_csv('heart_disease_dataset.csv')
```

```
#print first 5 rows
heart_data.head()
```



	Age	Gender	Cholesterol	Blood Pressure	Heart Rate	Smoking	Alcohol Intake	Exercise Hours	Family History	Diagnosis
0	75	Female	228	119	66	Current	Heavy	1	No	
1	48	Male	204	165	62	Current	NaN	5	No	
2	53	Male	234	91	67	Never	Heavy	3	Yes	
3	69	Female	192	90	72	Current	NaN	4	No	
4	62	Female	172	163	93	Never	NaN	6	No	

```
#print last 5 rows
heart_data.tail()
```



	Age	Gender	Cholesterol	Blood Pressure	Heart Rate	Smoking	Alcohol Intake	Exercise Hours	Family History
995	56	Female	269	111	86	Never	Heavy	5	No
996	78	Female	334	145	76	Never	NaN	6	No
997	79	Male	151	179	81	Never	Moderate	4	Yes
998	60	Female	326	151	68	Former	NaN	8	Yes
999	53	Male	226	116	82	Current	NaN	6	No

```
# no of rows and columns in the dataset
heart_data.shape
```



```
(1000, 16)
```

```
# getting some info about the data
heart_data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1000 non-null   int64
1   Gender                               1000 non-null   object
2   Cholesterol                           1000 non-null   int64
3   Blood Pressure                        1000 non-null   int64
4   Heart Rate                            1000 non-null   int64
5   Smoking                               1000 non-null   object
6   Alcohol Intake                        660 non-null    object
7   Exercise Hours                        1000 non-null   int64
8   Family History                        1000 non-null   object
9   Diabetes                              1000 non-null   object
10  Obesity                               1000 non-null   object
11  Stress Level                          1000 non-null   int64
12  Blood Sugar                           1000 non-null   int64
13  Exercise Induced Angina               1000 non-null   object
14  Chest Pain Type                       1000 non-null   object
15  Heart Disease                         1000 non-null   int64
dtypes: int64(8), object(8)
memory usage: 125.1+ KB
```

```
# checking for missing values
heart_data.isnull().sum()
```



	0
Age	0
Gender	0
Cholesterol	0
Blood Pressure	0
Heart Rate	0
Smoking	0
Alcohol Intake	340
Exercise Hours	0
Family History	0
Diabetes	0
Obesity	0
Stress Level	0
Blood Sugar	0
Exercise Induced Angina	0
Chest Pain Type	0
Heart Disease	0

dtype: int64

```
# statistical measures about the data
heart_data.describe()
```



	Age	Cholesterol	Blood Pressure	Heart Rate	Exercise Hours	Stress Level	
count	1000.000000	1000.000000	1000.0000	1000.000000	1000.000000	1000.000000	1000
mean	52.293000	249.939000	135.2810	79.204000	4.529000	5.646000	134
std	15.727126	57.914673	26.3883	11.486092	2.934241	2.831024	36
min	25.000000	150.000000	90.0000	60.000000	0.000000	1.000000	70
25%	39.000000	200.000000	112.7500	70.000000	2.000000	3.000000	104
50%	52.000000	248.000000	136.0000	79.000000	4.500000	6.000000	135
75%	66.000000	299.000000	159.0000	89.000000	7.000000	8.000000	167
max	79.000000	349.000000	179.0000	99.000000	9.000000	10.000000	199

```
# checking the distribution of the target value
heart_data['Heart Disease'].value_counts()
```



	count
Heart Disease	
0	608
1	392

dtype: int64

prompt: check for duplicate values from the code

```
from google.colab import drive
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Importing the Dependencies
drive.mount('/content/drive')

# Data collection and Processing
# loading the csv data to Pandas
heart_data = pd.read_csv('heart_disease_dataset.csv')

# print first 5 rows
print(heart_data.head())

# print last 5 rows
print(heart_data.tail())

# number of rows and columns in the dataset
print(heart_data.shape)

# getting some info about the data
print(heart_data.info())

# checking for missing values
print(heart_data.isnull().sum())

# statistical measures about the data
print(heart_data.describe())

# checking the distribution of the target value
print(heart_data['Heart Disease'].value_counts())

# Check for duplicate rows
duplicate_rows = heart_data[heart_data.duplicated()]
print("Number of duplicate rows:", len(duplicate_rows))

# Display duplicate rows (optional)
duplicate_rows
```

```
# Remove duplicate rows (optional - uncomment to remove)
# heart_data = heart_data.drop_duplicates()
# print("Number of rows after removing duplicates:", len(heart_data))
```

➡ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive

	Age	Gender	Cholesterol	Blood Pressure	Heart Rate	Smoking	\
0	75	Female	228	119	66	Current	
1	48	Male	204	165	62	Current	
2	53	Male	234	91	67	Never	
3	69	Female	192	90	72	Current	
4	62	Female	172	163	93	Never	

	Alcohol Intake	Exercise Hours	Family History	Diabetes	Obesity	\
0	Heavy	1	No	No	Yes	
1	NaN	5	No	No	No	
2	Heavy	3	Yes	No	Yes	
3	NaN	4	No	Yes	No	
4	NaN	6	No	Yes	No	

	Stress Level	Blood Sugar	Exercise Induced Angina	Chest Pain Type	\
0	8	119	Yes	Atypical Angina	
1	9	70	Yes	Typical Angina	
2	5	196	Yes	Atypical Angina	
3	7	107	Yes	Non-anginal Pain	
4	2	183	Yes	Asymptomatic	

	Heart Disease
0	1
1	0
2	1
3	0
4	0

	Age	Gender	Cholesterol	Blood Pressure	Heart Rate	Smoking	\
995	56	Female	269	111	86	Never	
996	78	Female	334	145	76	Never	
997	79	Male	151	179	81	Never	
998	60	Female	326	151	68	Former	
999	53	Male	226	116	82	Current	

	Alcohol Intake	Exercise Hours	Family History	Diabetes	Obesity	\
995	Heavy	5	No	Yes	Yes	
996	NaN	6	No	No	No	
997	Moderate	4	Yes	No	Yes	
998	NaN	8	Yes	Yes	No	
999	NaN	6	No	No	Yes	

	Stress Level	Blood Sugar	Exercise Induced Angina	Chest Pain Type	\
995	10	120	No	Non-anginal Pain	
996	10	196	Yes	Typical Angina	
997	8	189	Yes	Asymptomatic	
998	5	174	Yes	Atypical Angina	
999	5	161	Yes	Asymptomatic	

	Heart Disease
995	1
996	1
997	0
998	1
999	1

(1000, 16)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

```

---
0   Age               1000 non-null  int64
1   Gender            1000 non-null  object
2   Cholesterol       1000 non-null  int64
3   Blood Pressure    1000 non-null  int64
4   Heart Rate        1000 non-null  int64
5   Smoking           1000 non-null  object
6   Alcohol Intake    660 non-null   object
7   Exercise Hours    1000 non-null  int64
8   Family History    1000 non-null  object
9   Diabetes          1000 non-null  object
10  Obesity           1000 non-null  object
11  Stress Level      1000 non-null  int64
12  Blood Sugar       1000 non-null  int64
13  Exercise Induced Angina 1000 non-null  object
14  Chest Pain Type   1000 non-null  object
15  Heart Disease     1000 non-null  int64

```

dtypes: int64(8), object(8)

memory usage: 125.1+ KB

None

```

Age               0
Gender            0
Cholesterol       0
Blood Pressure    0
Heart Rate        0
Smoking           0
Alcohol Intake    340
Exercise Hours    0
Family History    0
Diabetes          0
Obesity           0
Stress Level      0
Blood Sugar       0
Exercise Induced Angina 0
Chest Pain Type   0
Heart Disease     0

```

dtype: int64

	Age	Cholesterol	Blood Pressure	Heart Rate	Exercise Hours \
count	1000.000000	1000.000000	1000.0000	1000.000000	1000.000000
mean	52.293000	249.939000	135.2810	79.204000	4.529000
std	15.727126	57.914673	26.3883	11.486092	2.934241
min	25.000000	150.000000	90.0000	60.000000	0.000000
25%	39.000000	200.000000	112.7500	70.000000	2.000000
50%	52.000000	248.000000	136.0000	79.000000	4.500000
75%	66.000000	299.000000	159.0000	89.000000	7.000000
max	79.000000	349.000000	179.0000	99.000000	9.000000

	Stress Level	Blood Sugar	Heart Disease
count	1000.000000	1000.000000	1000.000000
mean	5.646000	134.941000	0.392000
std	2.831024	36.699624	0.488441
min	1.000000	70.000000	0.000000
25%	3.000000	104.000000	0.000000
50%	6.000000	135.000000	0.000000
75%	8.000000	167.000000	1.000000
max	10.000000	199.000000	1.000000

Heart Disease

0 608

1 392

Name: count, dtype: int64

```
Number of duplicate rows: 0
```


Age	Gender	Cholesterol	Blood Pressure	Heart Rate	Smoking	Alcohol Intake	Exercise Hours	Family History


```
# prompt: create barchart for age in the dataset
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Assuming 'heart_data' DataFrame is already loaded as in your provided code
```

```
plt.figure(figsize=(8, 6))
sns.countplot(x='Age', data=heart_data)
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



```
# prompt: create barchart for sex in dataset
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```


```
# Assuming 'heart_data' DataFrame is already loaded as in your provided code
```

```
# Assuming 'Sex' column exists. If not, replace 'Sex' with the actual column name.
# Changed 'Sex' to 'Gender' based on the global variables
plt.figure(figsize=(6, 4))
```

```

sns.countplot(x='Gender', data=heart_data) # Using 'Gender' instead of 'Sex'
plt.title('Distribution of Sex')
plt.xlabel('Sex')
plt.ylabel('Count')
plt.show()

```



```

import matplotlib.pyplot as plt
import seaborn as sns

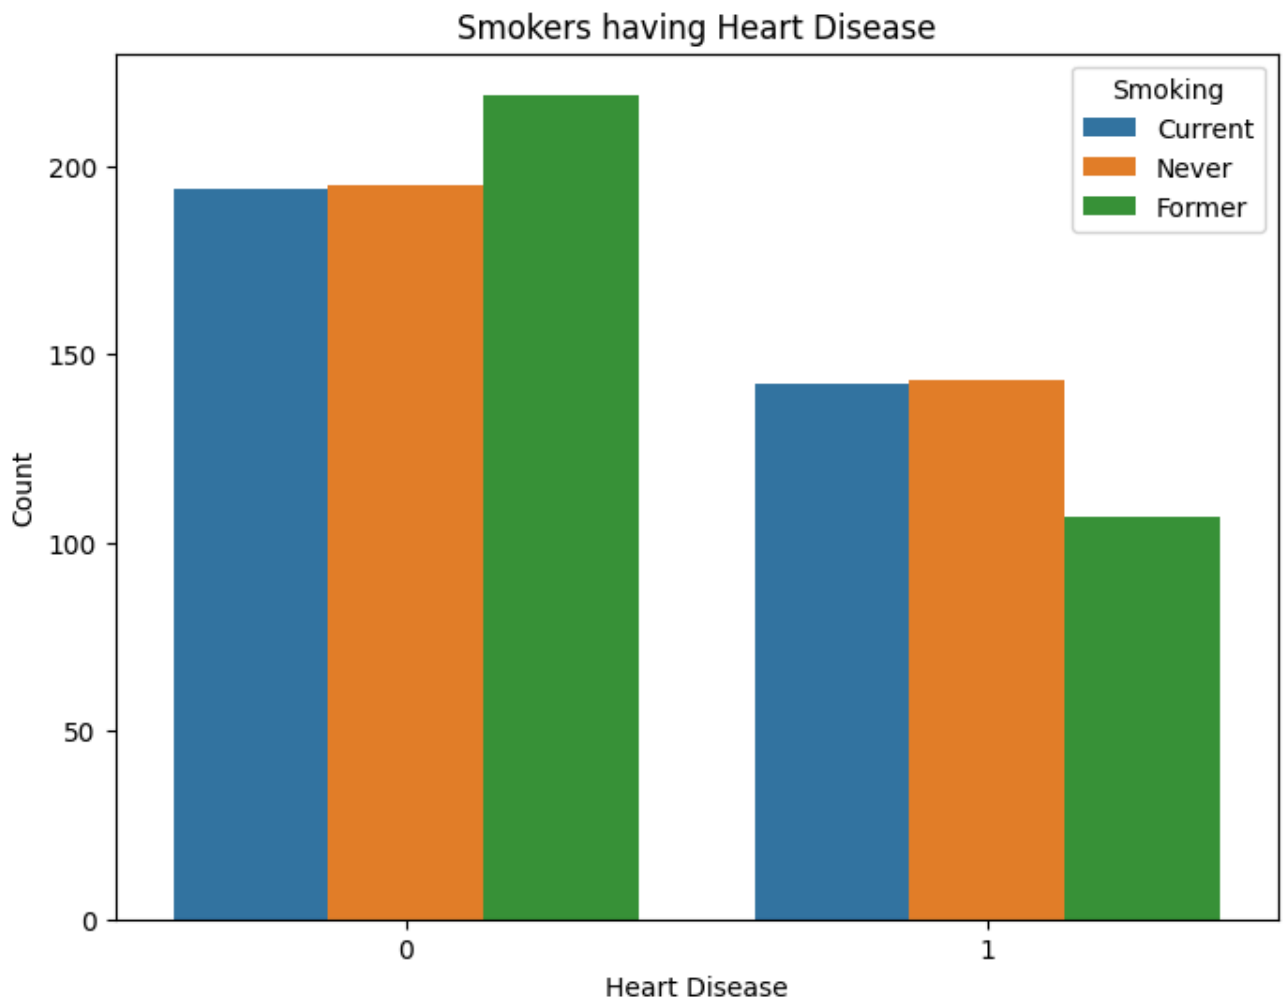
```

Assuming 'heart_data' DataFrame is already loaded as in your provided code

```

plt.figure(figsize=(8, 6))
# Replace 'Smoker' with the actual column name in your dataset, for example, 'Smoking'
sns.countplot(x='Heart Disease', hue='Smoking', data=heart_data)
plt.title('Smokers having Heart Disease')
plt.xlabel('Heart Disease')
plt.ylabel('Count')
plt.show()

```



✓ Bar chart for visualizing the person does not have heart disease and person having heart disease


prompt: create an bar chart for visualizing the person does not have heart disease |

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Assuming 'heart_data' DataFrame is already loaded as in your previous code
```

```
# Count the occurrences of each target value
heart_disease_counts = heart_data['Heart Disease'].value_counts()
```

```
# Create a bar chart
plt.figure(figsize=(8, 6))
sns.barplot(x=heart_disease_counts.index, y=heart_disease_counts.values)
plt.title('Distribution of Heart Disease')
plt.xlabel('Heart Disease (0 = No, 1 = Yes)')
plt.ylabel('Number of Patients')
plt.show()
```



- ✓ create an pie chart for visualizing the person does not have heart disease and person having heart disease

```
# prompt: create an pie chart for visualizing the person does not have heart disease


# Assuming 'heart_data' DataFrame is already loaded as in your previous code

# Count the occurrences of each target value
heart_disease_counts = heart_data['Heart Disease'].value_counts()

# Create a pie chart
plt.figure(figsize=(8, 6))
plt.pie(heart_disease_counts, labels=heart_disease_counts.index, autopct='%1.1f%%', shadow=True)
plt.title('Distribution of Heart Disease')
plt.show()
```



Distribution of Heart Disease



Start coding or [generate](#) with AI.

1--> Defective heart 0--> Healthy heart

Splitting the features and data

```
X = heart_data.drop(columns='Heart Disease', axis=1)
Y = heart_data['Heart Disease']
```

```
print(X)
```



	Age	Gender	Cholesterol	Blood Pressure	Heart Rate	Smoking	\
0	75	Female	228	119	66	Current	
1	48	Male	204	165	62	Current	
2	53	Male	234	91	67	Never	
3	69	Female	192	90	72	Current	
4	62	Female	172	163	93	Never	
..	
995	56	Female	269	111	86	Never	
996	78	Female	334	145	76	Never	
997	79	Male	151	179	81	Never	
998	60	Female	326	151	68	Former	
999	53	Male	226	116	82	Current	

	Alcohol Intake	Exercise Hours	Family History	Diabetes	Obesity	\
0	Heavy	1	No	No	Yes	
1	NaN	5	No	No	No	
2	Heavy	3	Yes	No	Yes	
3	NaN	4	No	Yes	No	
4	NaN	6	No	Yes	No	
..	
995	Heavy	5	No	Yes	Yes	
996	NaN	6	No	No	No	
997	Moderate	4	Yes	No	Yes	
998	NaN	8	Yes	Yes	No	
999	NaN	6	No	No	Yes	

	Stress Level	Blood Sugar	Exercise Induced Angina	Chest Pain Type
0	8	119	Yes	Atypical Angina
1	9	70	Yes	Typical Angina
2	5	196	Yes	Atypical Angina
3	7	107	Yes	Non-anginal Pain
4	2	183	Yes	Asymptomatic
..
995	10	120	No	Non-anginal Pain
996	10	196	Yes	Typical Angina
997	8	189	Yes	Asymptomatic
998	5	174	Yes	Atypical Angina
999	5	161	Yes	Asymptomatic

[1000 rows x 15 columns]

```
print(Y)
```

```

0      1
1      0
2      1
3      0
4      0
..
995    1
996    1
997    0
998    1
999    1
Name: Heart Disease, Length: 1000, dtype: int64
```

prompt: create an heatmap for the given data

Importing necessary libraries

```
import pandas as pd # Import pandas for data manipulation
```

```
import matplotlib.pyplot as plt # Import matplotlib for plotting
```

```
import seaborn as sns # Import seaborn for heatmap visualization
```

Assuming 'heart_data' DataFrame is already loaded as in your previous code

and X, Y are defined as in the previous code.

Load the dataset (if not already loaded)

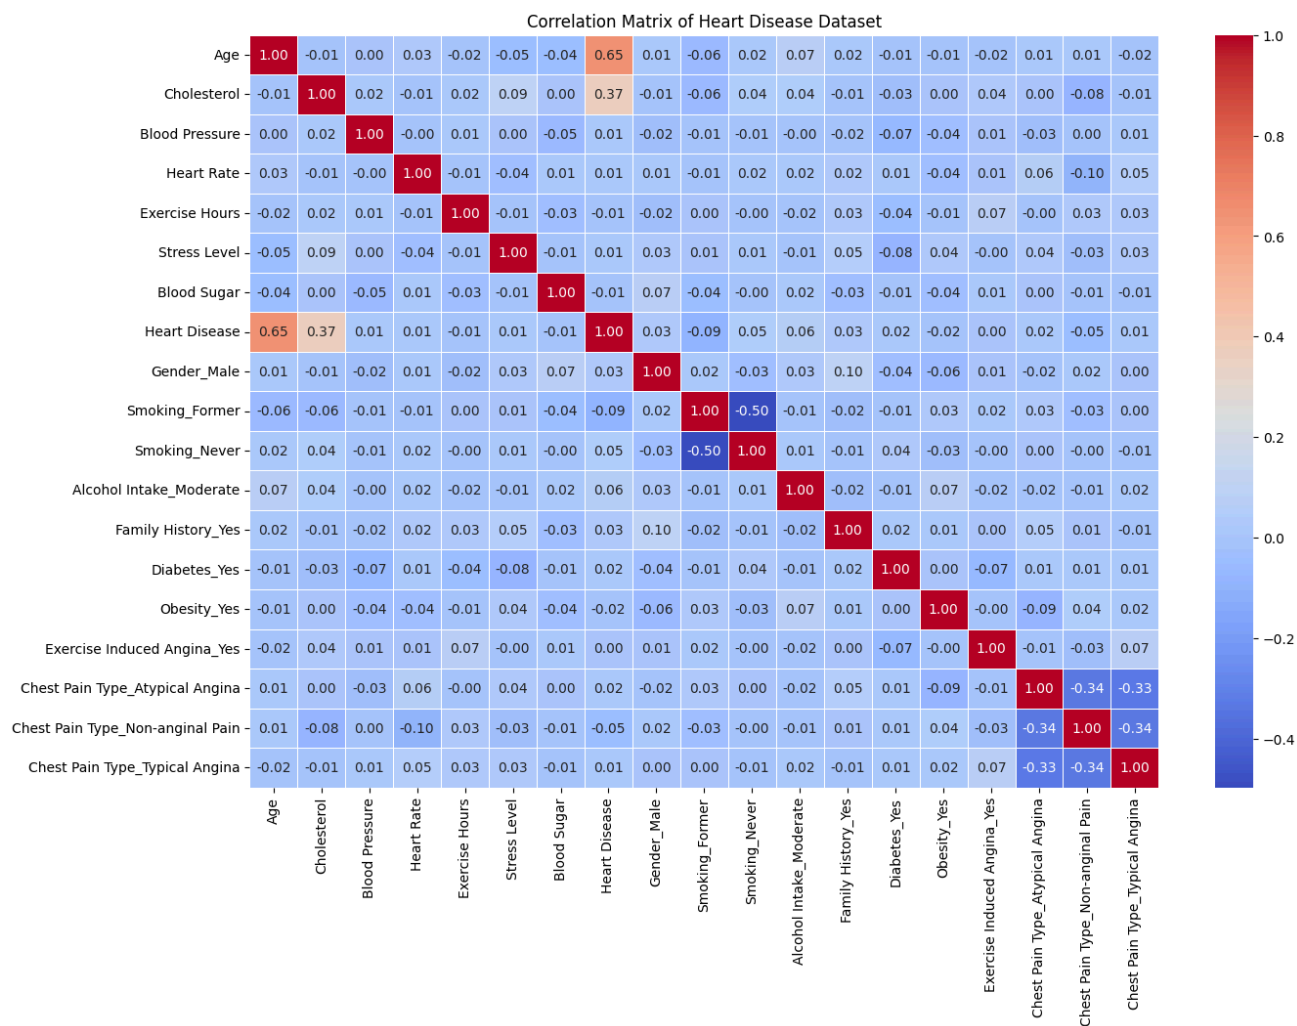
```
heart_data = pd.read_csv('heart_disease_dataset.csv') # Load the data into heart_dat
```

Convert categorical columns to numerical using one-hot encoding

```
# Select categorical features (object type)
categorical_features = heart_data.select_dtypes(include=['object']).columns
# Perform one-hot encoding
heart_data_encoded = pd.get_dummies(heart_data, columns=categorical_features, drop_first=True)

# Calculate the correlation matrix using the encoded data
correlation_matrix = heart_data_encoded.corr()

# Create a heatmap
plt.figure(figsize=(15, 10)) # Adjust figure size as needed
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Correlation Matrix of Heart Disease Dataset")
plt.show()
```



Splitting the data into Training data and Testing data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, rand
```

```
print(X.shape, X_train.shape, X_test.shape)
```



```
(1000, 15) (800, 15) (200, 15)
```

Model Training

Logistic Regression

```
model = LogisticRegression()
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
# Assuming 'heart_data' is your original DataFrame

# ... (Your previous code for loading and inspecting the data) ...

# Splitting the features and data
X = heart_data.drop(columns='Heart Disease', axis=1)
Y = heart_data['Heart Disease']

# Get the actual column names from the DataFrame
print(X.columns)

# Perform one-hot encoding on categorical features
# Update the column names to reflect the actual column names
categorical_features = X.select_dtypes(include=['object']).columns.tolist() #Selects
# Alternatively, manually adjust the column names to match those printed above
# categorical_features = ['Sex', 'ChestPainType', 'FastingBS', 'RestingECG', 'Exercis
print(categorical_features)
X = pd.get_dummies(X, columns=categorical_features, drop_first=True) # drop_first to

# Splitting the data into Training data and Testing data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y,

# Model Training
model = LogisticRegression()

# training the logistic regression model with training data
model.fit(X_train, Y_train)
```



```

➡ Index(['Age', 'Gender', 'Cholesterol', 'Blood Pressure', 'Heart Rate',
        'Smoking', 'Alcohol Intake', 'Exercise Hours', 'Family History',
        'Diabetes', 'Obesity', 'Stress Level', 'Blood Sugar',
        'Exercise Induced Angina', 'Chest Pain Type'],
        dtype='object')
['Gender', 'Smoking', 'Alcohol Intake', 'Family History', 'Diabetes', 'Obesity', 'Exe
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: Conver
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
    ▾ LogisticRegression ⓘ ?
```

```
LogisticRegression()
```

Model Evaluation

✓ Accuracy Score

```

# accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)

```

```
print('Accuracy on Training data : ', training_data_accuracy)
```

```

# accuracy on training data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)

```

```
print('Accuracy on Test data : ', test_data_accuracy)
```

```
➡ Accuracy on Test data : 0.845
```

```
# prompt: create confusion create for all the methods used in the program
```

```

from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

```

```
# ... (Your existing code) ...
```

```
# Model Evaluation
```

```
# ### Accuracy Score
```

```
# accuracy on training data
```

```
X_train_prediction = model.predict(X_train)
```

```
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```



```
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print('Accuracy on Training data : ', training_data_accuracy)
```


```
# accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print('Accuracy on Test data : ', test_data_accuracy)
```

```
# Confusion Matrix for Training Data
cm_train = confusion_matrix(Y_train, X_train_prediction)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_train, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title('Confusion Matrix for Training Data')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

```
# Confusion Matrix for Test Data
cm_test = confusion_matrix(Y_test, X_test_prediction)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_test, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title('Confusion Matrix for Test Data')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

⇒ Accuracy on Training data : 0.8425
Accuracy on Test data : 0.83





Predicted 0

Predicted 1

Predicted Label

✓ Building Predictive System

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Assuming 'heart_data' is your original DataFrame and you've loaded it

# ... (Your previous code for loading and inspecting the data) ...

# Get the actual column names from the DataFrame
# ... (Your code for splitting and printing X.columns) ...

# Perform one-hot encoding on categorical features during training
# ... (Your code for one-hot encoding during training) ...

### Building Predictive System

input_data = (50, 'Male', 308, 166, 97, 'Current', 'None', 7, 'No', 'No', 'No', 2, 116, 'Yes', 'Typical')

# Convert input_data to DataFrame for one-hot encoding
# creates a single row dataframe with the same columns as the ORIGINAL heart_data
input_data_df = pd.DataFrame([input_data], columns=heart_data.columns[:-1])
#print(input_data_df) # check what the columns are

# Perform one-hot encoding on the input data using the same columns as in training data
input_data_encoded = pd.get_dummies(input_data_df, columns=categorical_features, drop_first=True)

# Ensure that the input data has the same columns as the training data after encoding
missing_cols = set(X_train.columns) - set(input_data_encoded.columns) #checks if any columns are missing
for c in missing_cols:
    input_data_encoded[c] = 0 #Add any missing column and set its value to 0 to ensure consistency
input_data_encoded = input_data_encoded[X_train.columns] # Reorder columns to match training data

# Convert the encoded DataFrame back to a NumPy array for prediction
input_data_resaped = input_data_encoded.to_numpy()

prediction = model.predict(input_data_resaped)
print(prediction)

if (prediction[0] == 0):
    print('The Person does not have a Heart Disease')
```

```
else:
    print('The Person has Heart Disease')

[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning
    warnings.warn(
```

✓ SVM METHOD

```
# prompt: produce prediction by using support vector machine

from sklearn.svm import SVC

# ... (Your existing code for data loading, preprocessing, and one-hot encoding) ...

# Model Training using Support Vector Machine
model = SVC(kernel='linear') # You can experiment with different kernels (e.g., 'rbf', '
model.fit(X_train, Y_train)

# Model Evaluation
# ... (Your existing code for model evaluation using accuracy_score) ...

# Building Predictive System
# ... (Your existing code for creating input_data and encoding it) ...

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')
```

```
[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning
    warnings.warn(
```

✓ Accuracy for this svm method

```
# prompt: accuracy for this svm method

# ... (Your existing code for data loading, preprocessing, and one-hot encoding) ...

# Model Training using Support Vector Machine
model = SVC(kernel='linear') # You can experiment with different kernels (e.g., 'rbf', '
model.fit(X_train, Y_train)
```

```

# Model Evaluation
# Accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print('Accuracy on Training data : ', training_data_accuracy)

# Accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print('Accuracy on Test data : ', test_data_accuracy)
# Building Predictive System
# ... (Your existing code for creating input_data and encoding it) ...

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')

```

```

➡ Accuracy on Training data : 0.85875
Accuracy on Test data : 0.885
[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning
    warnings.warn(

```

✓ create prediction in random forest method with prediction accuracy

```

# prompt: create prediction in random forest method with prediction accuracy

from sklearn.ensemble import RandomForestClassifier

# ... (Your existing code for data loading, preprocessing, and one-hot encoding) ...

# Model Training using Random Forest
model = RandomForestClassifier(random_state=42) # You can adjust hyperparameters
model.fit(X_train, Y_train)

# Model Evaluation
# Accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print('Accuracy on Training data : ', training_data_accuracy)

# Accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print('Accuracy on Test data : ', test_data_accuracy)

```

```
# Building Predictive System
# ... (Your existing code for creating input_data and encoding it) ...
```

```
prediction = model.predict(input_data_reshaped)
print(prediction)
```

```
if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')
```

```
➡ Accuracy on Training data : 1.0
Accuracy on Test data : 1.0
[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning
    warnings.warn(
```

```
# prompt: create accuracy,precision, recall, r2 score for the random forest method train
```

```
# Import necessary libraries
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import precision_score, recall_score, r2_score, accuracy_score
import pandas as pd
from sklearn.model_selection import train_test_split
```

```
# ... (Your existing code for data loading, preprocessing, and model training) ...
# Assuming 'heart_data' is your original DataFrame and you've loaded it
```

```
# ... (Your previous code for loading and inspecting the data) ...
```

```
# Splitting the features and data
X = heart_data.drop(columns='Heart Disease', axis=1)
Y = heart_data['Heart Disease']
```

```
# Perform one-hot encoding on categorical features
categorical_features = X.select_dtypes(include=['object']).columns.tolist()
X = pd.get_dummies(X, columns=categorical_features, drop_first=True)
```

```
# Splitting the data into Training data and Testing data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, rand
```

```
# Model Training using Random Forest
model = RandomForestClassifier(random_state=42)
model.fit(X_train, Y_train)
```

```
# Predictions
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)
```

```
# Evaluate the model
print("Random Forest Performance:")
```

```

print("Train Set:")
print(f"Accuracy: {accuracy_score(Y_train, y_train_pred)}")
print(f"Precision: {precision_score(Y_train, y_train_pred)}")
print(f"Recall: {recall_score(Y_train, y_train_pred)}")
print(f"R2 Score: {r2_score(Y_train, y_train_pred)}") # Note: R2 score might not be ideal

print("\nTest Set:")
print(f"Accuracy: {accuracy_score(Y_test, y_test_pred)}")
print(f"Precision: {precision_score(Y_test, y_test_pred)}")
print(f"Recall: {recall_score(Y_test, y_test_pred)}")
print(f"R2 Score: {r2_score(Y_test, y_test_pred)}") # Note: R2 score might not be ideal f

➡ Random Forest Performance:
Train Set:
Accuracy: 1.0
Precision: 1.0
Recall: 1.0
R2 Score: 1.0

Test Set:
Accuracy: 1.0
Precision: 1.0
Recall: 1.0
R2 Score: 1.0

```

✓ create prediction in knn method with prediction accuracy

```

# prompt: create prediction in knn method with prediction accuracy

from sklearn.neighbors import KNeighborsClassifier

# ... (Your existing code for data loading, preprocessing, and one-hot encoding) ...

# Model Training using K-Nearest Neighbors
model = KNeighborsClassifier(n_neighbors=5) # You can adjust the number of neighbors
model.fit(X_train, Y_train)

# Model Evaluation
# Accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print('Accuracy on Training data : ', training_data_accuracy)

# Accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print('Accuracy on Test data : ', test_data_accuracy)

# Building Predictive System
# ... (Your existing code for creating input_data and encoding it) ...

prediction = model.predict(input_data_reshaped)
print(prediction)

```



```

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')

```

```

➡ Accuracy on Training data : 0.95125
Accuracy on Test data : 0.92
[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning
    warnings.warn(

```

✓ create prediction in gradient boosting method with prediction accuracy

```

# prompt: create prediction in gradient boosting method with prediction accuracy

from sklearn.ensemble import GradientBoostingClassifier

# ... (Your existing code for data loading, preprocessing, and one-hot encoding) ...

# Model Training using Gradient Boosting
model = GradientBoostingClassifier(random_state=42) # You can adjust hyperparameters
model.fit(X_train, Y_train)

# Model Evaluation
# Accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print('Accuracy on Training data : ', training_data_accuracy)

# Accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print('Accuracy on Test data : ', test_data_accuracy)

# Building Predictive System
# ... (Your existing code for creating input_data and encoding it) ...

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')

```

```

➡ Accuracy on Training data : 1.0
Accuracy on Test data : 1.0
[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning

```

```
warnings.warn(
```

✓ create prediction in naive bayes method with prediction accuracy

```
# prompt: create prediction in naive bayes method with prediction accuracy

from sklearn.naive_bayes import GaussianNB

# ... (Your existing code for data loading, preprocessing, and one-hot encoding) ...

# Model Training using Naive Bayes
model = GaussianNB()
model.fit(X_train, Y_train)

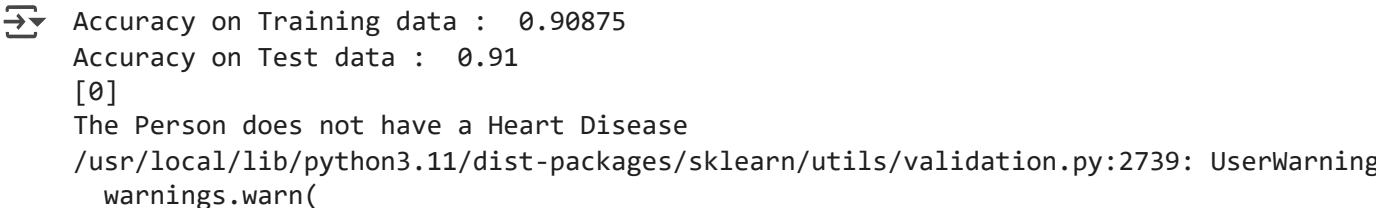
# Model Evaluation
# Accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print('Accuracy on Training data : ', training_data_accuracy)

# Accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print('Accuracy on Test data : ', test_data_accuracy)

# Building Predictive System
# ... (Your existing code for creating input_data and encoding it) ...

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')
```



```
⇒ Accuracy on Training data : 0.90875
Accuracy on Test data : 0.91
[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning
warnings.warn(
```

✓ create prediction in clustering method with prediction accuracy

```
# prompt: create prediction in clustering method with prediction accuracy

from sklearn.cluster import KMeans
```

```

from sklearn.metrics import accuracy_score

# Assuming X_train, X_test, Y_train, Y_test are already defined from your previous code

# Model Training using KMeans (Clustering)
kmeans = KMeans(n_clusters=2, random_state=42) # Assuming 2 clusters for binary classifi
kmeans.fit(X_train)

# Predict cluster labels for the training and testing data
X_train_cluster_labels = kmeans.predict(X_train)
X_test_cluster_labels = kmeans.predict(X_test)

# Map cluster labels to target labels (0 or 1) based on majority class in each cluster
cluster_mapping = {} # Dictionary to store the mapping
for cluster_label in range(2):
    cluster_data = Y_train[X_train_cluster_labels == cluster_label]
    majority_class = cluster_data.mode()[0]
    cluster_mapping[cluster_label] = majority_class

# Convert predicted cluster labels to target labels using the mapping
Y_train_predicted = np.array([cluster_mapping[cluster_label] for cluster_label in X_train
Y_test_predicted = np.array([cluster_mapping[cluster_label] for cluster_label in X_test_c

# Model Evaluation
training_accuracy = accuracy_score(Y_train, Y_train_predicted)
testing_accuracy = accuracy_score(Y_test, Y_test_predicted)

print('Accuracy on Training data : ', training_accuracy)
print('Accuracy on Test data : ', testing_accuracy)

# Building Predictive System (for the same input_data as before)
input_data_cluster_label = kmeans.predict(input_data_reshaped)
predicted_target_label = cluster_mapping[input_data_cluster_label[0]]

print(predicted_target_label)

if predicted_target_label == 0:
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')

```



```

Accuracy on Training data : 0.635
Accuracy on Test data : 0.675
1
The Person has Heart Disease
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning
warnings.warn(

```

```

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

```